

# Class 7: Machine Learning I

Yujia Liu (PID:A16967405)

Today we are going to learn how to apply different machine learning methods, beginning with clustering:

The goal here is to find cluster/groups in your input data.


First I will make up some data with clear groups. For this I will use the `rnorm()` function.

```
rnorm(10) #Give 10 random numbers from normal distribution
```


```
[1] 0.38942456 -1.28573724 0.49745097 1.60649675 -0.02975181 -0.22809807  
[7] -1.30972086 0.13125049 0.37959831 1.84418140
```

```
hist(rnorm(10000, mean = 3))
```

**Histogram of `rnorm(10000, mean = 3)`**



```
n <- 10000
x <- c(rnorm(n,-3), rnorm(n), +3) #Make a vector of normal distribution
hist(x)
```




```
n <- 30
x <- c(rnorm(n,-3), rnorm(n), +3)
y <- rev(x)

z <- cbind(x,y)
head(z)
```

	x	y
[1,]	-4.2386906	3.0000000
[2,]	-0.6970244	0.2811774
[3,]	-3.7866359	-1.6358185
[4,]	-2.9739680	-2.6444166
[5,]	-3.2327851	-0.8898777
[6,]	-3.4127389	-0.4318250

```
plot(z)
```



Use the `kmeans()` function setting `k` to 2 and `nstart=20`

Inspect/print the results

Q. How many points are in each cluster?

Q. What 'component' of your result object details - cluster size? - cluster assignment/membership? - cluster center?

```
km <- kmeans(z, centers = 2)
km
```

K-means clustering with 2 clusters of sizes 31, 30

Cluster means:

	x	y
1	-0.4232245	-3.0860075
2	-3.1805521	-0.4290097

Clustering vector:

```
[1] 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 1 1 1 1 1 1 1 1 1
[39] 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1
```


Within cluster sum of squares by cluster:




Q. Plot x colored by the kmeans cluster assignment and add cluster centers as blue points

R will re-cycle the shorter color vector to be the same length as the longer (number of data points) in z

```
plot(z, col=c("red", "blue"))
```




```
plot(z, col = km$cluster)
points(km$centers, col = "blue", pch = 15, cex=3) #Make mean shown on plot.
```



Q. Can you run kmeans and ask for 4 clusters please and plot the results like we have done above?

```
km4 <- kmeans(z,centers = 4)
plot(z, col = km4$cluster)
points(km4$centers, col = "blue",pch = 15, cex=3)
```



```
#It will be different every time you run it.
```

##Hierarchical Clustering

Let's take our same made-up data **z** and see how **hclust** works.

First we need a distance matrix of our data to be clustered.

```
d <- dist(z)
hc <- hclust(d)
hc
```


Call:

```
hclust(d = d)
```

```
Cluster method   : complete
Distance          : euclidean
Number of objects: 61
```

```
plot(hc)
abline(h=8, col="red")
```

## Cluster Dendrogram



```
hclust (*, "complete")
```

I can get my cluster membership vector by “cutting the tree” with the `cutree()` function like so:


```
grps <- cutree(hc, h=8)
grps
```

```
[1] 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 2 2 2 2 2 2 2
[39] 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 1 2 1 2
```

Can you plot **z** colored by our hclust results:

```
plot(z, col = grps)
```






##PCA of UK food data

Read data from the UK on food consumption in different parts of the UK.

```
url <- "https://tinyurl.com/UK-foods"
x <- read.csv(url, row.names=1)
head(x)
```

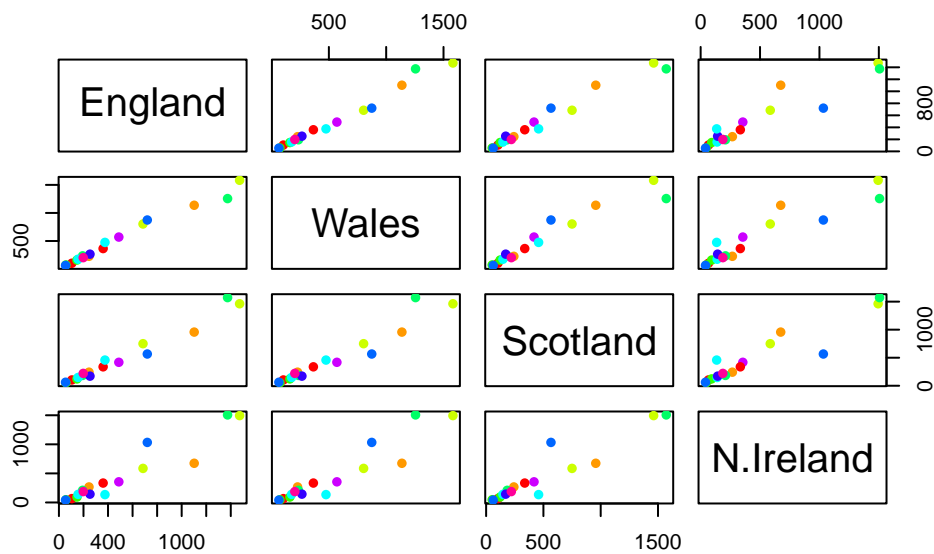
	England	Wales	Scotland	N.Ireland
Cheese	105	103	103	66
Carcass_meat	245	227	242	267
Other_meat	685	803	750	586
Fish	147	160	122	93
Fats_and_oils	193	235	184	209
Sugars	156	175	147	139

```
barplot(as.matrix(x), beside=T, col=rainbow(nrow(x)))
```



A so-called “Pairs” plot can be useful for small datasets like this:

```
pairs(x, col=rainbow(10), pch=16)
```



It is hard to see structure and trends in even this small data-set. How will we ever do this when we have big datasets with 10002 or 10s of thousands of things we are measuring...

###PCA to the rescue

Let's see how PCA deals with this dataset. So main function in base R to do PCA is called `prcomp()`

```
pca <- prcomp(t(x))
summary(pca)
```

Importance of components:

	PC1	PC2	PC3	PC4
Standard deviation	324.1502	212.7478	73.87622	2.921e-14
Proportion of Variance	0.6744	0.2905	0.03503	0.000e+00
Cumulative Proportion	0.6744	0.9650	1.00000	1.000e+00

Let's see what's inside this `pca` object that we created from running `prcomp()`

```
attributes(pca)
```

\$names

```
[1] "sdev"      "rotation" "center"    "scale"     "x"
```

\$class

```
[1] "prcomp"
```

```
pca$x
```

	PC1	PC2	PC3	PC4
England	-144.99315	-2.532999	105.768945	-9.152022e-15
Wales	-240.52915	-224.646925	-56.475555	5.560040e-13
Scotland	-91.86934	286.081786	-44.415495	-6.638419e-13
N.Ireland	477.39164	-58.901862	-4.877895	1.329771e-13

```
plot(pca$x[,1], pca$x[,2],
     col = c("black", "red", "blue", "darkgreen"),
     pch=16,
     xlab="PC1 (67.4%)",
     ylab="PC2 (29%)")
```

